

HARIPRASAD BATHINI SANKARAN

📍 Cincinnati, Ohio 📞 +1 (513) 227-1668 ✉️ hariprasad.sankaran@gmail.com 🌐 hariprasadbs 🎧 Hariprasad-b-s 📁 Portfolio

MASTER'S STUDENT | SEEKING DATA ENGINEERING INTERNSHIP (SUMMER 2026)

SUMMARY

Master's student in Information Technology with hands-on experience in data engineering. Experienced in building scalable data pipelines using **Python**, **SQL**, **Spark**, **Databricks**, **AWS**, and **Azure**. Seeking a Data Engineering Internship for Summer 2026 to apply industry experience in a research-driven and production environment.

EDUCATION

University of Cincinnati

Master of Science in Information Technology

Aug 2025 - Dec 2026

Cincinnati, Ohio, USA

Anna University

Bachelor of Engineering in Electronics and Communication Engineering

Jun 2016 – Nov 2020

TamilNadu, India

SKILLS

Data Engineering: Databricks, Snowflake, Apache Spark, PySpark, SQL, Apache Airflow, dbt, ETL/ELT, Data Warehousing, SQL Server, Data Modeling, Delta Lake, Medallion Architecture, Streaming Data, Data Quality REST API, Apache Kafka, Unity Catalog

Programming: Python (Pandas, NumPy, Scikit-learn), SQL, PySpark

Cloud Platforms: AWS (S3, Lambda, Glue, Redshift), Azure (ADLS Gen2, Data Factory, SQL Server), Terraform

Tools: Git, GitHub, VS Code, Cribl, Jupyter, Unix/Linux, Agile,

PROJECTS

Security Log Analytics and Threat Detection Platform (Azure)

[Link](#) 

- Built large-scale data pipelines using **Azure Data Factory** and **Databricks** to collect and process security telemetry from simulated sources.
- Implemented transformation workflows using **PySpark** and stored curated datasets in **ADLS Gen2** using Parquet and Delta formats.
- Designed Bronze–Silver–Gold **medallion architecture** to support real-time analytics and historical reporting use cases.
- Implemented data validation checks and alerting logic to evaluate pipeline reliability and **data quality** in a production-like environment.
- Designed operational-style dashboards to monitor ingestion latency, throughput, and system health for proactive troubleshooting.

Real-Time Retail Demand Forecasting Pipeline (AWS)

- Designed and implemented an automated ETL pipeline using **AWS S3**, **Glue**, **Lambda**, and **Redshift** to process 50K+ retail events/hour.
- Built data ingestion and transformation workflows using **Python** and **PySpark**, handling structured and semi-structured datasets at scale.
- Optimized SQL queries and partition strategies, reducing dashboard refresh latency and improving query performance by **30%**.
- Developed interactive dashboards in **Power BI** to track inventory trends, improving demand forecasting accuracy by **25%**.

WORK EXPERIENCE

Data Engineer | Comcast Corporation | Chennai, India

Aug 2023 – Aug 2025

- Built scalable batch and streaming data pipelines in **Databricks** and **Snowflake** to process high-volume cybersecurity telemetry logs for threat analytics.
- Developed orchestration workflows using **Apache Airflow** to automate ingestion, transformation, and validation across multi-stage pipelines.
- Implemented **structured streaming** pipelines to enable near real-time detection of anomalous security events and system alerts.
- Designed reusable transformation modules and dbt models to standardize data processing across multiple data sources.
- Managed infrastructure as code using **Terraform** and implemented version control for configuration files using **Git**.
- Optimized storage layout and compute performance using **Delta Lake OPTIMIZE** and **ZORDER**, reducing annual cloud costs by **\$20K**.
- Collaborated with analysts and security engineers to translate detection requirements into scalable data solutions and reporting datasets.
- Created reusable ingestion templates and Airflow DAG generators, reducing onboarding time for new data sources by **40%**.

Data Engineer | Infosys Ltd. (Microsoft Corporation) | Chennai, India

Jun 2021 – Aug 2023

- Built and optimized ETL pipelines in **Azure Databricks** integrating raw operational datasets into curated analytical layers in **ADLS Gen2**.
- Converted complex T-SQL stored procedures into **PySpark** jobs, reducing batch processing time by **30 minutes** per run.
- Designed and implemented metadata-driven ingestion frameworks in **Azure Data Factory** to support scalable onboarding of new data.
- Developed REST API ingestion frameworks in Python with retry logic, pagination handling, and schema validation, accelerating pipeline development by **50%**.
- Built multi-dimensional data models (fact and dimension tables) in **Databricks** to support BI dashboards and performance reporting.
- Established automated data quality checks using Python to detect schema drift and prevent corrupted downstream analytics.

CERTIFICATION

• **Databricks Certified Data Engineer Professional**

Dec 2025

• **Databricks LakeHouse Fundamentals**

Feb 2024

• **Azure Data Fundamentals (DP-900)**

Mar 2022

• **Azure Fundamentals (AZ-900)**

Jul 2021